

CISC 5950 - Big Data Programming

Project 2 - Part 2 Report

Team Members:

Revathi Bhuvaneswari

Tianying Luo

Youfei Zhang

Yue Zeng

Spring 2019

P2: Heart Disease Prediction using Logistic Regression

Bash script: Import the data file from HDFS.

```
#!/bin/bash
#../start.sh
source ../env.sh
/usr/local/hadoop/bin/hdfs dfs -rm -r /part2/input/
/usr/local/hadoop/bin/hdfs dfs -rm -r /part2/output/
/usr/local/hadoop/bin/hdfs dfs -mkdir -p /part2/input/
/usr/local/hadoop/bin/hdfs dfs -copyFromLocal ../data/framingham.csv /part2/input/
/usr/local/spark/bin/spark-submit --master=spark://$SPARK_MASTER:7077 ./framingham.py hdfs://$SPARK_MASTER:9000/part2/input/
#../stop.sh
```

Python Script:

Stage 1: The train / test data are read from the CSV files and loaded into dataframes.

```
reload(sys)
sys.setdefaultencoding('utf8')
if __name__ == "__main__":
    if len(sys.argv) != 2:
        print("Usage: Q2.py <directory_of_files>", file=sys.stderr)
        sys.exit(-1)

    spark = SparkSession.builder.appName('lr-predic').getOrCreate()
    df = spark.read.csv(sys.argv[1], header = True, inferSchema = True)
    df.show()
    cols=df.columns
    df.printSchema()
```

Stage 2: The dataset is cleaned, the vectors of 'features' & 'label' are added to the dataframe.

```
#Filter out all the 'NA' values
df = df.filter((df.cigsPerDay != 'NA') & (df.BPMeds != 'NA') & (df.totChol != 'NA') & (df.BMI != 'NA') & (df.heartRate != 'NA') & (df.glucose != 'NA'))
string_cols=['cigsPerDay', 'BPMeds', 'totChol', 'BMI', 'heartRate', 'glucose']
for col_name in string_cols:
    df = df.withColumn(col_name, col(col_name).cast('float'))

#Count of the class values
df.groupby('TenYearCHD').count().show()

#Drop the column of 'education'
#Generative statistics conclusion
cols.remove('education')
cols.insert(0, 'Summary')
df.describe().select(cols[:len(cols)//2]).show()
df.describe().select(cols[len(cols)//2:-1]).show()

#Use pipeline to combine all the features in one single feature vector
cols.remove('TenYearCHD')
cols.remove('Summary')
print(cols)

['male', 'age', 'currentSmoker', 'cigsPerDay', 'BPMeds', 'prevalentStroke', 'prevalentHyp',
 'diabetes', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose']

df = ohe(df)
df.printSchema()
```

The cleaning process is done by first filtering out all the 'NA' values from the following columns: 'cigsPerDay', 'BPMeds', 'totChol', 'BMI', 'heartRate', 'glucose'. 'NA' values represent missing values in the dataset and are identified as string value. Then all the columns of string type are altered by converting them into float type. Additionally, the unwanted column of 'education' is dropped.

The following function is used to one-hot-encode all categorical columns, and get vector representation of all features. String Indexer and One Hot Encoder Estimator are used for this.

```

# converts sparse vector to dense vector
def sparse_to_array(v):
    v = DenseVector(v)
    new_array = list([float(x) for x in v])
    return new_array

# returns one-hot-encoded stages
def ohe(df):
    # One-hot-encode category columns
    categoricalColumns = ['education', 'currentSmoker', 'BPMeds', 'prevalentStroke', 'prevalentHyp', 'diabetes']
    stages = []
    for categoricalCol in categoricalColumns:
        stringIndexer = StringIndexer(inputCol = categoricalCol, outputCol = categoricalCol + 'Index')
        encoder = OneHotEncoderEstimator(inputCols=[stringIndexer.getOutputCol()], outputCols=[categoricalCol + "classVec"])
        stages += [stringIndexer, encoder]

    # vectorize numeric columns
    numericCols = ['male', 'age', 'cigsPerDay', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose']
    assemblerInputs = [c + "classVec" for c in categoricalColumns] + numericCols
    assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
    stages += [assembler]

    # perform final pipeline to get vectorized feature column
    pipeline = Pipeline().setStages(stages)
    pipelineModel = pipeline.fit(df)
    df = pipelineModel.transform(df)
    selectedCols = ['TenYearCHD', 'features'] + cols
    df = df.select(selectedCols)

    # convert sparse vector initial_feature column to dense vector
    sparse_to_array_udf = udf(sparse_to_array, T.ArrayType(T.FloatType()))
    df = df.withColumn('dense_vector_features', sparse_to_array_udf('features'))
    ud_f = udf(lambda r : Vectors.dense(r), VectorUDT())
    df = df.withColumn('features', ud_f('dense_vector_features'))
    return df

```

Stage 3: In this stage, a new column of ‘Scaled_features’ is created by using StandardScaler to standardize the newly feature values according to their different calibration. Also, since the classification values are imbalanced. Balancing ratio is calculated, based on which class weights will be calculated and applied to different classes. In addition, ChiSqSelector method is used to select the feature with p-value below 0.05, which proves a strong predictive power. At last, use the class column as labelCol, ‘Aspect’ (generated by chi-square feature selection) as featuresCol, and classweights as weightCol, and the iteration time is set to be 10. After that, split the dataset into train and test by 80/20. Train and fit the model by using the training dataset.

```

#Standardization
standardscaler=StandardScaler().setInputCol("features").setOutputCol("Scaled_features")
df=standardscaler.fit(df).transform(df)
df.select("features", "Scaled_features").show(5)

#Split train & test
train, test = df.randomSplit([0.8, 0.2], seed=123)

total=float(train.select("TenYearCHD").count())
numPositives=train.select("TenYearCHD").where('TenYearCHD == 1').count()
per_ones=(float(numPositives)/float(total))*100
numNegatives=float(total-numPositives)
print('\n\nThe number of Class 1 are {}'.format(numPositives))
print('\n\nPercentage of Class 1 are {}'.format(per_ones))

#Add balancing ratio
BalancingRatio= numNegatives/total
print('\n\nBalancingRatio = {}'.format(BalancingRatio))

#Creat a new column named "classWeights" in the "train" dataset
train=train.withColumn("classWeights", when(train.TenYearCHD==1, BalancingRatio).otherwise(1-BalancingRatio))
train.select("classWeights", "TenYearCHD").show(10)

#Feature selection
css = ChiSqSelector(featuresCol='features', outputCol='Aspect', labelCol='TenYearCHD', fpr=0.05)
train=css.fit(train).transform(train)
test=css.fit(test).transform(test)
test.select("Aspect").show(5, truncate=False)

#Train the lr model
lr = LogisticRegression(labelCol="TenYearCHD", featuresCol="Aspect", weightCol="classWeights", maxIter=10)
model=lr.fit(train)

```

Stage 4: In this last stage, the trained Logistic Regression model is used to make predictions by applying the transform() function on both the train as well as the test data frames. Then extract the counts of TP, TN, FP, and FN by filtering the corresponding pairs and counting them. The Area Under the ROC is calculated using the Binary Classification Evaluator by passing in the actual predictions data frame, which has both the labels and predictions.

```
predict_train=model.transform(train)
predict_test=model.transform(test)
predict_train.select("TenYearCHD","prediction").show(10)
predict_test.select("TenYearCHD","prediction").show(10)

#Evaluation of the train model
TP_train = predict_train.filter((predict_train.TenYearCHD == 1) & (predict_train.prediction == 1.0)).count()
TN_train = predict_train.filter((predict_train.TenYearCHD == 0) & (predict_train.prediction == 0.0)).count()
FP_train = predict_train.filter((predict_train.TenYearCHD == 0) & (predict_train.prediction == 1.0)).count()
FN_train = predict_train.filter((predict_train.TenYearCHD == 1) & (predict_train.prediction == 0.0)).count()
evaluation(TP_train, TN_train, FP_train, FN_train)

#Evaluation of the test model
TP_test = predict_test.filter((predict_test.TenYearCHD == 1) & (predict_test.prediction == 1.0)).count()
TN_test = predict_test.filter((predict_test.TenYearCHD == 0) & (predict_test.prediction == 0.0)).count()
FP_test = predict_test.filter((predict_test.TenYearCHD == 0) & (predict_test.prediction == 1.0)).count()
FN_test = predict_test.filter((predict_test.TenYearCHD == 1) & (predict_test.prediction == 0.0)).count()
evaluation(TP_test, TN_test, FP_test, FN_test)

#ROC for train & test
evaluator=BinaryClassificationEvaluator(rawPredictionCol="rawPrediction",labelCol="TenYearCHD")
predict_test.select("TenYearCHD","rawPrediction","prediction","probability").show(5, truncate = False)

print("\n\nThe area under ROC for train set is {}".format(evaluator.evaluate(predict_train)))
print("\n\nThe area under ROC for test set is {}".format(evaluator.evaluate(predict_test)))
```

The following function is used to calculate and output the accuracy, misclassification rate, sensitivity, specificity, positive/negative predictive value and positive/negative likelihood ratio.

```
def evaluation(TP, TN, FP, FN):
    print('\n\n', 'The confusion matrix is: \n\n', '[[' ,TN,FP,']\n\n', '[[' ,FN,TP,']]]', '\n\n')
    #Model Evaluation - Statistics
    sensitivity=TP/float(TP+FN)
    specificity=TN/float(TN+FP)

    print('\n\nThe accuracy of the model = (TP+TN)/(TP+TN+FP+FN) = ',(TP+TN)/float(TP+TN+FP+FN),'\n\n',
    'The Missclassification = 1-Accuracy = ',1-((TP+TN)/float(TP+TN+FP+FN)),'\n\n',
    'Sensitivity or True Positive Rate = TP/(TP+FN) = ',TP/float(TP+FN),'\n\n',
    'Specificity or True Negative Rate = TN/(TN+FP) = ',TN/float(TN+FP),'\n\n',
    'Positive Predictive value = TP/(TP+FP) = ',TP/float(TP+FP),'\n\n',
    'Negative predictive Value = TN/(TN+FN) = ',TN/float(TN+FN),'\n\n',
    'Positive Likelihood Ratio = Sensitivity/(1-Specificity) = ',sensitivity/(1-specificity),'\n\n',
    'Negative likelihood Ratio = (1-Sensitivity)/Specificity = ',(1-sensitivity)/specificity,'\n\n')
```

Output:

male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	TenYearCHD
1	39	4	0	0	0	0	0	0	195	106.0	70.0	26.97	80	77	0
0	46	2	0	0	0	0	0	0	250	121.0	81.0	28.73	95	76	0
1	48	1	1	20	0	0	0	0	245	127.5	80.0	25.34	75	70	0
0	61	3	1	30	0	0	1	0	225	150.0	95.0	28.58	65	103	1
0	46	3	1	23	0	0	0	0	285	130.0	84.0	23.1	85	85	0
0	43	2	0	0	0	0	1	0	228	180.0	110.0	30.3	77	99	0
0	63	1	0	0	0	0	0	0	205	138.0	71.0	33.11	60	85	1
0	45	2	1	20	0	0	0	0	313	100.0	71.0	21.68	79	78	0
1	52	1	0	0	0	0	1	0	260	141.5	89.0	26.36	76	79	0
1	43	1	1	30	0	0	1	0	225	162.0	107.0	23.61	93	88	0
0	50	1	0	0	0	0	0	0	254	133.0	76.0	22.91	75	76	0
0	43	2	0	0	0	0	0	0	247	131.0	88.0	27.64	72	61	0
1	46	1	1	15	0	0	1	0	294	142.0	94.0	26.31	98	64	0
0	41	3	0	0	1	0	1	0	332	124.0	88.0	31.31	65	84	0
0	39	2	1	9	0	0	0	0	226	114.0	64.0	22.35	85	NA	0
0	38	2	1	20	0	0	1	0	221	140.0	90.0	21.35	95	70	1
1	48	3	1	10	0	0	1	0	232	138.0	90.0	22.37	64	72	0
0	46	2	1	20	0	0	0	0	291	112.0	78.0	23.38	80	89	1
0	38	2	1	5	0	0	0	0	195	122.0	84.5	23.24	75	78	0
1	41	2	0	0	0	0	0	0	195	139.0	88.0	26.88	85	65	0

Figure 1: Initial DataFrame (first 20 rows)

```

root
|-- male: integer (nullable = true)
|-- age: integer (nullable = true)
|-- education: string (nullable = true)
|-- currentSmoker: integer (nullable = true)
|-- cigsPerDay: string (nullable = true)
|-- BPMeds: string (nullable = true)
|-- prevalentStroke: integer (nullable = true)
|-- prevalentHyp: integer (nullable = true)
|-- diabetes: integer (nullable = true)
|-- totChol: string (nullable = true)
|-- sysBP: double (nullable = true)
|-- diaBP: double (nullable = true)
|-- BMI: string (nullable = true)
|-- heartRate: string (nullable = true)
|-- glucose: string (nullable = true)
|-- TenYearCHD: integer (nullable = true)

```

Figure 2: Schema of the Initial DataFrame

TenYearCHD	count
1	644
0	3596

TenYearCHD	count
1	572
0	3179

Figure 3: Classification Counts Before & After Data Cleaning

Summary	male	age	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
count	3751	3751	3751	3751	3751	3751	3751
mean	0.4452146094374833	49.57344708077846	0.48840309250866437	9.008531058384431	0.03039189549453479	0.005598507064782724	0.3119168221807518
stddev	0.49705575870704916	8.570204436756383	0.49993213793450775	11.925096527979754	0.17168601101128225	0.07462337676919495	0.4633378373062201
min	0	32	0	0.0	0.0	0	0
max	1	70	1	70.0	1.0	1	1

diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose
3751	3751	3751	3751	3751	3751	3751
0.027192748600373233	236.92801919488136	132.3684350839776	82.93854972007465	25.80828845268947	75.70407891229006	81.88003199146894
0.1626663986140648	44.61159360579599	22.04652246327564	11.932779343310061	4.06559867696019	11.956382146147071	23.882233365111478
0	113.0	83.5	48.0	15.54	44.0	40.0
1	696.0	295.0	142.5	56.8	143.0	394.0

Figure 4: Generative Statistics Analysis of Each Attribute

```

root
|-- TenYearCHD: integer (nullable = true)
|-- features: vector (nullable = true)
|-- male: integer (nullable = true)
|-- age: integer (nullable = true)
|-- currentSmoker: integer (nullable = true)
|-- cigsPerDay: float (nullable = true)
|-- BPMeds: float (nullable = true)
|-- prevalentStroke: integer (nullable = true)
|-- prevalentHyp: integer (nullable = true)
|-- diabetes: integer (nullable = true)
|-- totChol: float (nullable = true)
|-- sysBP: double (nullable = true)
|-- diaBP: double (nullable = true)
|-- BMI: float (nullable = true)
|-- heartRate: float (nullable = true)
|-- glucose: float (nullable = true)
|-- dense_vector_features: array (nullable = true)
|   |-- element: float (containsNull = true)

```

Figure 5: Schema of the DataFrame After Data Cleaning

features	Scaled_features
[0.0,0.0,0.0,1.0,...]	[0.0,0.0,0.0,3.16...]
[0.0,1.0,0.0,0.0,...]	[0.0,2.1956978139...]
[1.0,0.0,0.0,0.0,...]	[2.03538650844691...]
[0.0,0.0,1.0,0.0,...]	[0.0,0.0,2.713096...]
[0.0,0.0,1.0,0.0,...]	[0.0,0.0,2.713096...]

only showing top 5 rows

Figure 6: Scaled_features of Standardization

```

The number of Class 1 are 460

Percentage of Class 1 are 15.343562374916612

BalancingRatio = 0.8465643762508339

```

Figure 7: Percentage of Class 1 & Balancing Ratio

classWeights	TenYearCHD
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0
0.15343562374916608	0

only showing top 10 rows

Aspect
[0.0,0.0,0.0,0.0,0.0,0.0,1.0,1.0,0.0,1.0,0.0,50.0,15.0,253.0,132.0,84.5,27.959999084472656,88.0,73.0]
[0.0,0.0,0.0,0.0,0.0,0.0,1.0,1.0,0.0,1.0,0.0,51.0,15.0,240.0,180.0,107.5,25.329999923706055,68.0,80.0]
[0.0,0.0,0.0,0.0,0.0,0.0,1.0,1.0,0.0,1.0,0.0,57.0,20.0,246.0,160.0,92.5,30.739999771118164,84.0,76.0]
[0.0,0.0,0.0,0.0,0.0,0.0,1.0,1.0,0.0,1.0,1.0,67.0,25.0,221.0,144.0,84.0,24.920000076293945,72.0,73.0]
[0.0,0.0,0.0,0.0,0.0,0.0,1.0,1.0,1.0,1.0,0.0,37.0,20.0,223.0,115.0,72.0,22.709999084472656,76.0,63.0]

only showing top 5 rows

Figure 8: Classweights & Aspect

TenYearCHD	prediction
0	1.0
0	1.0
0	1.0
0	1.0
0	1.0
0	1.0
0	0.0
0	0.0
0	0.0
0	0.0

only showing top 10 rows

TenYearCHD	prediction
0	0.0
0	1.0
0	1.0
0	1.0
0	0.0
0	0.0
0	0.0
0	0.0
0	0.0
0	1.0

only showing top 10 rows

Figure 9: Train & Test Prediction Outputs

TenYearCHD	rawPrediction	prediction	probability
0	[-0.1567527896942406, 0.1567527896942406]	1.0	[0.46089184842530045, 0.5391081515746997]
0	[-1.79364229862992, 1.79364229862992]	1.0	[0.14262674731316835, 0.8573732526868316]
0	[-1.616040225550969, 1.616040225550969]	1.0	[0.16575169569904596, 0.8342483043009541]
0	[0.9793847307922181, -0.9793847307922181]	0.0	[0.7269861164242096, 0.2730138835757904]
0	[0.6310721597565216, -0.6310721597565216]	0.0	[0.6527325313989591, 0.34726746860104085]

only showing top 5 rows

Figure 10: Probability of Predictions

The area under ROC for train set is 0.7269039999999907

Evaluation Results for Training Dataset

The confusion matrix is:

```
[[ 1677 823 ]
```

```
[ 154 296 ]]
```

The accuracy of the model = $(TP+TN)/(TP+TN+FP+FN) = 0.6688135593220339$

The Missclassification = $1 - \text{Accuracy} = 0.3311864406779661$

Sensitivity or True Positive Rate = $TP/(TP+FN) = 0.6577777777777778$

Specificity or True Negative Rate = $TN/(TN+FP) = 0.6708$

Positive Predictive value = $TP/(TP+FP) = 0.2645218945487042$

Negative predictive Value = $TN/(TN+FN) = 0.9158929546695794$

Positive Likelihood Ratio = $\text{Sensitivity}/(1 - \text{Specificity}) = 1.9981098960442822$

Negative likelihood Ratio = $(1 - \text{Sensitivity})/\text{Specificity} = 0.5101702776121381$

Figure 11: Evaluation Results for Training Dataset

The area under ROC for test set is 0.774487554021107

Evaluation Results for Test Dataset

The confusion matrix is:

```
[[ 482 197 ]
```

```
[ 33 89 ]]
```

The accuracy of the model = $(TP+TN)/(TP+TN+FP+FN) = 0.7128589263420724$

The Missclassification = $1 - \text{Accuracy} = 0.2871410736579276$

Sensitivity or True Positive Rate = $TP/(TP+FN) = 0.7295081967213115$

Specificity or True Negative Rate = $TN/(TN+FP) = 0.7098674521354934$

Positive Predictive value = $TP/(TP+FP) = 0.3111888111888112$

Negative predictive Value = $TN/(TN+FN) = 0.9359223300970874$

Positive Likelihood Ratio = $\text{Sensitivity}/(1 - \text{Specificity}) = 2.514396271948074$

Negative likelihood Ratio = $(1 - \text{Sensitivity})/\text{Specificity} = 0.38104550710835994$

Figure 12: Evaluation Results for Testing Dataset